

Nathan Gillette

The basis for this project is to microbenchmark two concurrent data structures to measure their throughput. The two data structures I chose were the built in `java.concurrent.ConcurrentHashMap` and a custom `ReadWriteLock` based `HashTable`. The Custom `HashTable` is read and written to in a 80/20 ratio to use the best performance practice for `ReadWriteLock` based data structures. The mock use case is a Flight Manifest, where each flight has an ID as well as a passenger number. The data structure can be initialized with a certain number of premade flights, or can be initialized with no flights. For testing purposes, I started the data structures with 20 premade flights. I hard coded these flights because I didn't want to skew my test with io. There were some extra methods needed in the `MyBenchmark` Class for the `jmh` benchmark, specifically the `setup` and `teardown` methods. These were crucial to the `jmh` benchmark.

I originally was considering other mock use cases such as a grocery store, a music and book library, and some kind of card game. I eventually settled on the Flight Manifest because I found some ideas online and it seemed like it would be straightforward enough while still working for the project requirements. It should come as no surprise that `jmh` was the hardest part of this project. Luckily, I was able to find a nice guide that uses maven to set up a nice template to work in. It's become abundantly clear to me that learning a language is really only half the battle, the tools and other infrastructure that is built around a language like Java is just as important.

Links:

[ScreenCap](#)

[Graph](#)